

Nano-RK:

an Energy-aware Resource-centric RTOS for Sensor Networks



Anand Eswaran, Anthony Rowe, Raj Rajkumar
{aeswaran, agr, raj}@ece.cmu.edu
Real-Time and Multimedia Systems Lab
Carnegie Mellon University

Presented by Anthony Rowe
RTSS 2005



Secure Sensor Networks for Physical Infrastructures

- Develop a secure **software platform for infrastructure sensor networks** to be deployed in physical structures such as factories, buildings, homes, bridges, campuses, vehicles, ships and planes.
- Provide **continual monitoring of operational health** and safety,
- **Report malfunctions** (nearly) instantly, and
- Support **mobile nodes** and **track people**.



What is nanoRK?

- **A Real-time Operating System for sensor nodes for use in wireless sensor networks**
 - **Priority-Based Preemptive Multitasking**
 - **Resource Reservations**
 - **Built-in Multi-hop Networking Support**



Related Work

- **Tiny OS (Berkeley)**
 - Large public following and support
 - Compact
 - Cooperative Multitasking
 - No timing guarantees for tasks
- **Mantis OS (University of Colorado)**
 - No real-time scheduling
 - No reservation support
- **uCOS, Emerald, OSEK**
 - Still too large
 - Not network centric

NanoRK Motivation

- **Quicker Development Cycles**
 - Traditional OS abstractions allow for quicker learning curves and help abstract away low level details
- **Scaling Technology**
 - Sensor Nodes will become increasingly complex as they begin to manage multiple tasks
 - Local Processing is much cheaper than using the network
- **Resource Reservations (Energy Budget)**
 - Reservations can enforce energy and communication budgets to minimize negative impact on network lifetimes from unintentional errors or malicious behavior on some nodes

NanoRK Motivation

- **Why Priority-based scheduling?**

	Period	Execution Time
Network Radio	Sporadic	10ms
Audio Sensor	200 hz	10us
Light Sensor	166 hz	10us
Smart Camera	1 hz	300ms
Global Positioning	5 hz	10ms

Time-triggered task interleaving can become daunting...

Furthermore, what if a new sensor is added or a period changes?

Goals of a Sensor RTOS (1 of 2)

- **Multitasking with Priority-Based Preemption**
 - Respond on a timely basis to critical events
 - Support predictability and schedulability
- **Built-in Network Support**
 - Decrease coding effort to implement custom communication protocols
- **Enforce Energy Usage Limits**
 - Meet Battery Lifetime Requirements

Goals of a Sensor RTOS (2 of 2)

- **Unified Sensor / Network API**
 - Sensors and actuators support in a device driver manner that is abstracted away from the user
- **Small Footprint**
 - Current trends in microcontrollers show larger ROM sizes (64Kb to 128KB) and smaller RAM sizes (2KB to 8KB)



FireFly Hardware

Atmega128L

CC2420 (802.15.4)

Light, Temp, Audio,
PIR, Acceleration,
Ultrasound

Global AM band
Synchronization



FireFly Synchronization Hardware



NanoRK Resources

Component	Resource
Context Swap Time	45 μ S
Mutex Structure Overhead	5 Bytes per Resource
Stack Size Per Task	32 \rightarrow 128 bytes (64 bytes by default)
OS Struct Overhead	50 bytes Per Task
Network Overhead	164 bytes
Total Typical Configuration: (8 tasks, 8 mutexes, 4 16 byte network buffers)	2KB RAM, 10KB ROM

Current Hardware Platform:


4Mhz Atmega128L with Chipcon CC2420 802.15.4 transceiver

NanoRK Reservations

- **CPU**
 - Each Task can be given a budget of how long it is allowed to execute per a given period
- **Network**
 - Per Task Budget on Network Usage
 - Transmit and Receive Packets and/or Bytes
- **Sensors / Actuators**
 - Number of System Calls to a particular peripheral per a given period

Virtual Energy Reservations

- **{CPU, Network, Sensors}**
 - Together comprise the total energy usage of the node
- **Static Offline Budget Enforcement**
 - It is possible to calculate a node lifetime given a certain energy budget

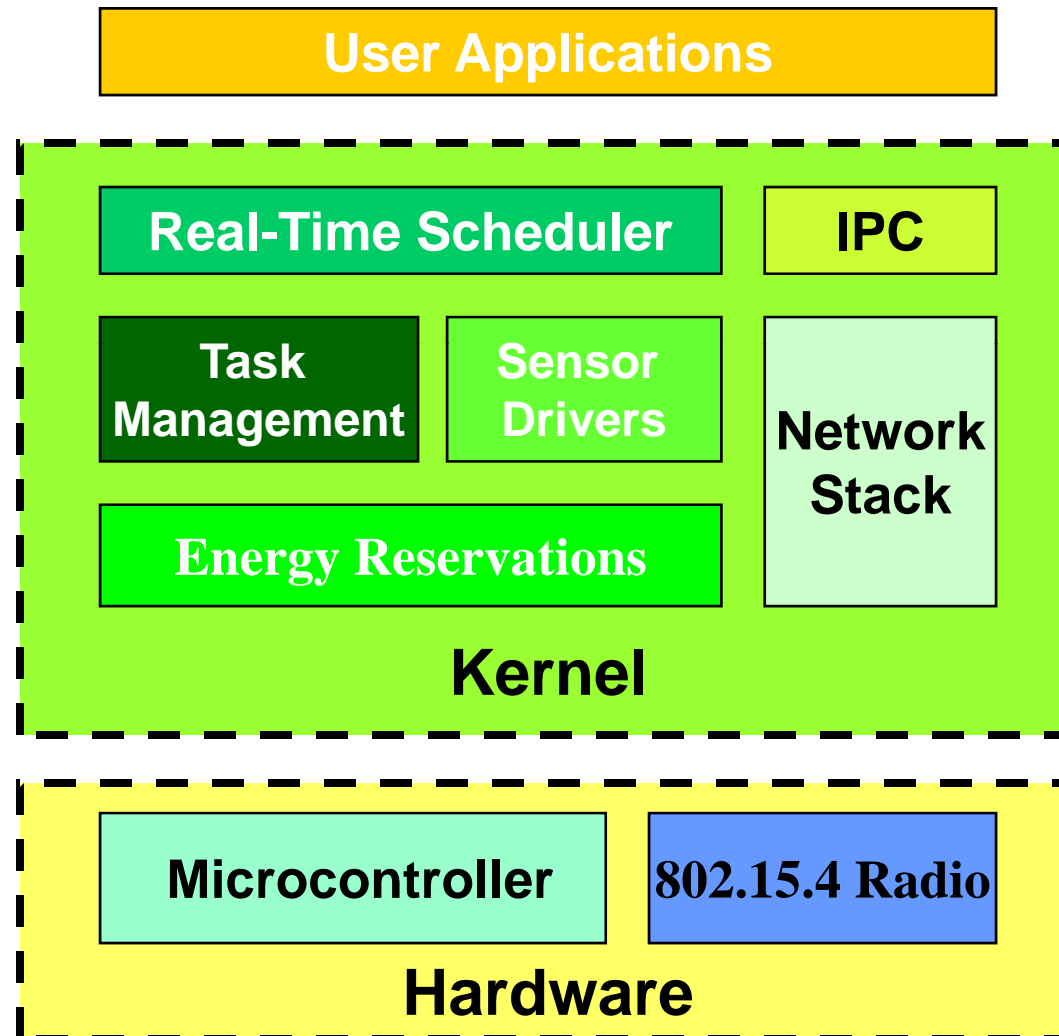


Node	Reserve [TX, RX]	TX Rate	Lifetime w/ out Reserve	Lifetime w/ Reserve
a	[1,2]	1	8 years	8 years
b	X	300	3.5 days	3.5 days
c	[1,2]	1	5 years	5 years
d	[1,2]	1	3.9 days	5 years
e	[1,2]	1	4 days	2.9 years

NanoRK Reservations

- **What if a reservation is violated?**
- **Hard Reservation**
 - For CPU reservations, the task is preempted
 - For Sensor/Network reservations, an error is returned to the task making the system call
- **Firm Reservations**
 - The task is allowed to consume slack resources until they have been depleted

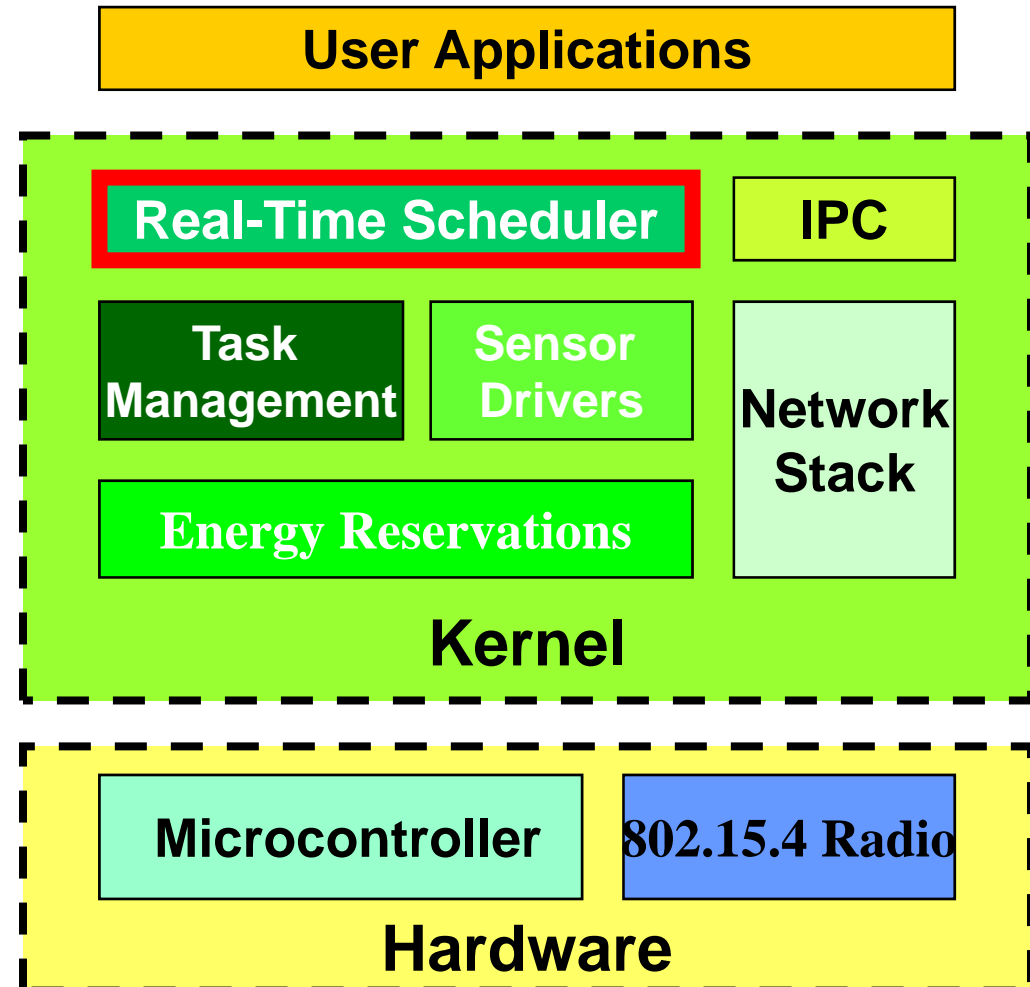
NanoRK Architecture



NanoRK Architecture

Real-Time Scheduler

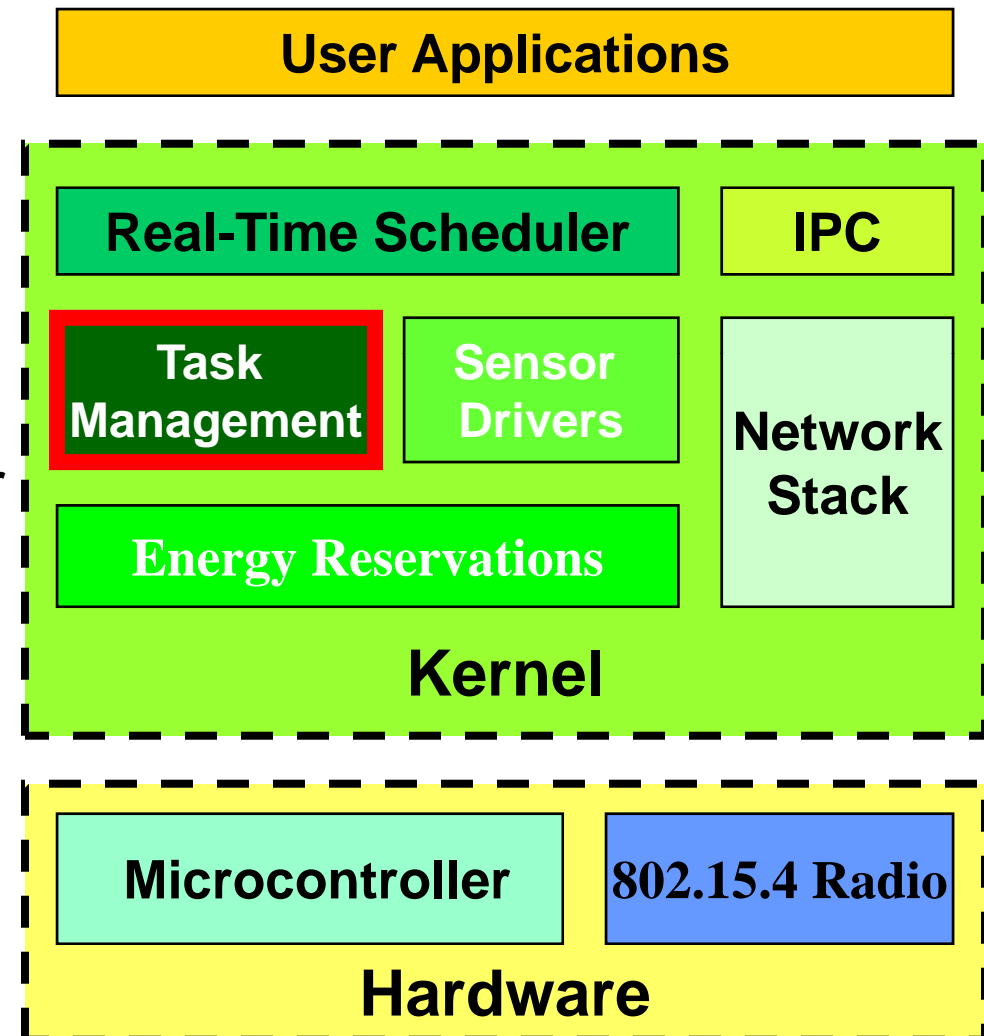
- Priority-based Scheduling
- Static Design-time Framework
- Priority Ceiling Protocol



NanoRK Architecture

- **Task Management**

- Two 32 bit counters
 - { Seconds, Nanoseconds }
- Periodic Task Switching
 - Triggered by a One Shot Timer or an Event
- Enables Fine Grained Timing Control



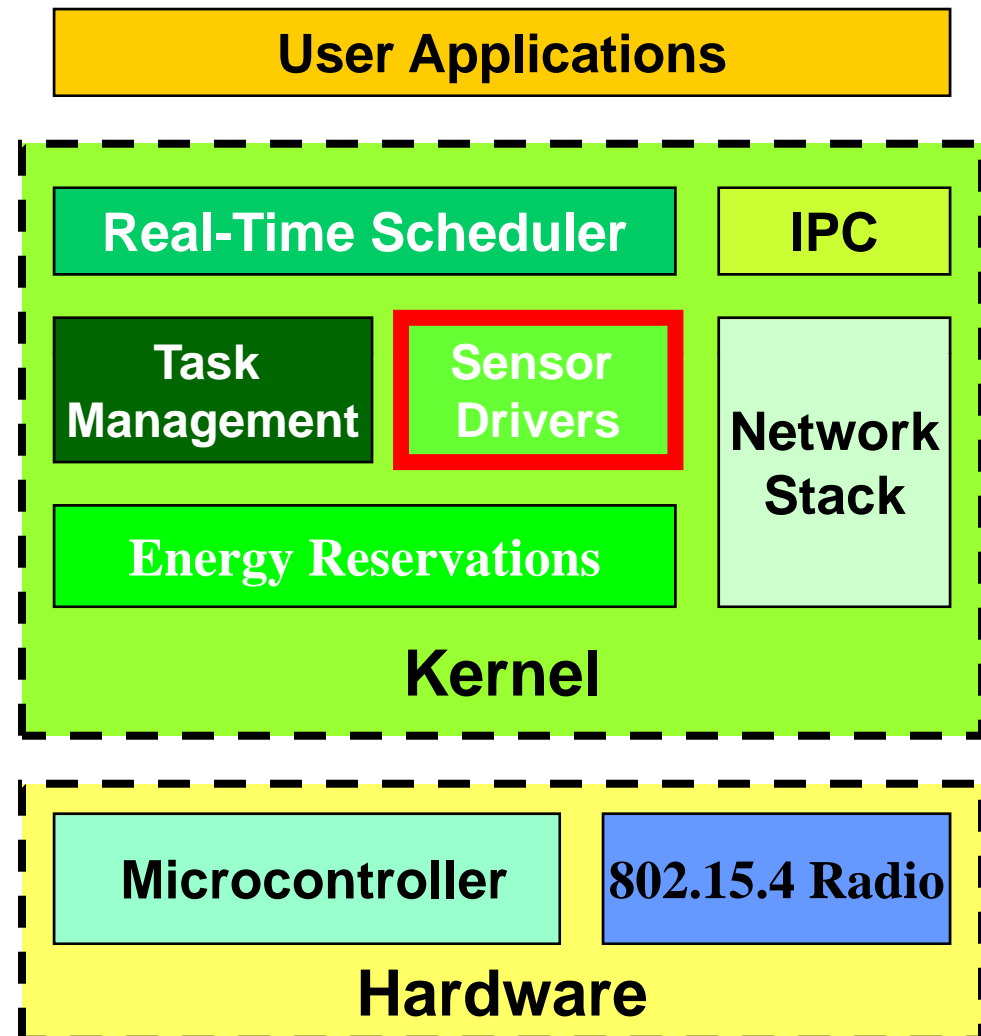
Nano-RK Task Specs

```
nrk_task_type Task1;  
  
Task1.task = Sound_Task;  
Task1.Ptos = (void *) &Stack1[STACKSIZE - 1];  
Task1.TaskID = 1;  
Task1.priority = 3;  
Task1.Period = 1000;    // ms  
Task1.set_reserve[CPU] = 50;  //ms  
Task1.set_reserve[NETWORK_PACKETS] = 3;  
Task1.set_reserve[SENSE_ACTUATE] = 12;  
nrk_activate_task(Task1);
```

NanoRK Architecture

- **Sensor Drivers**

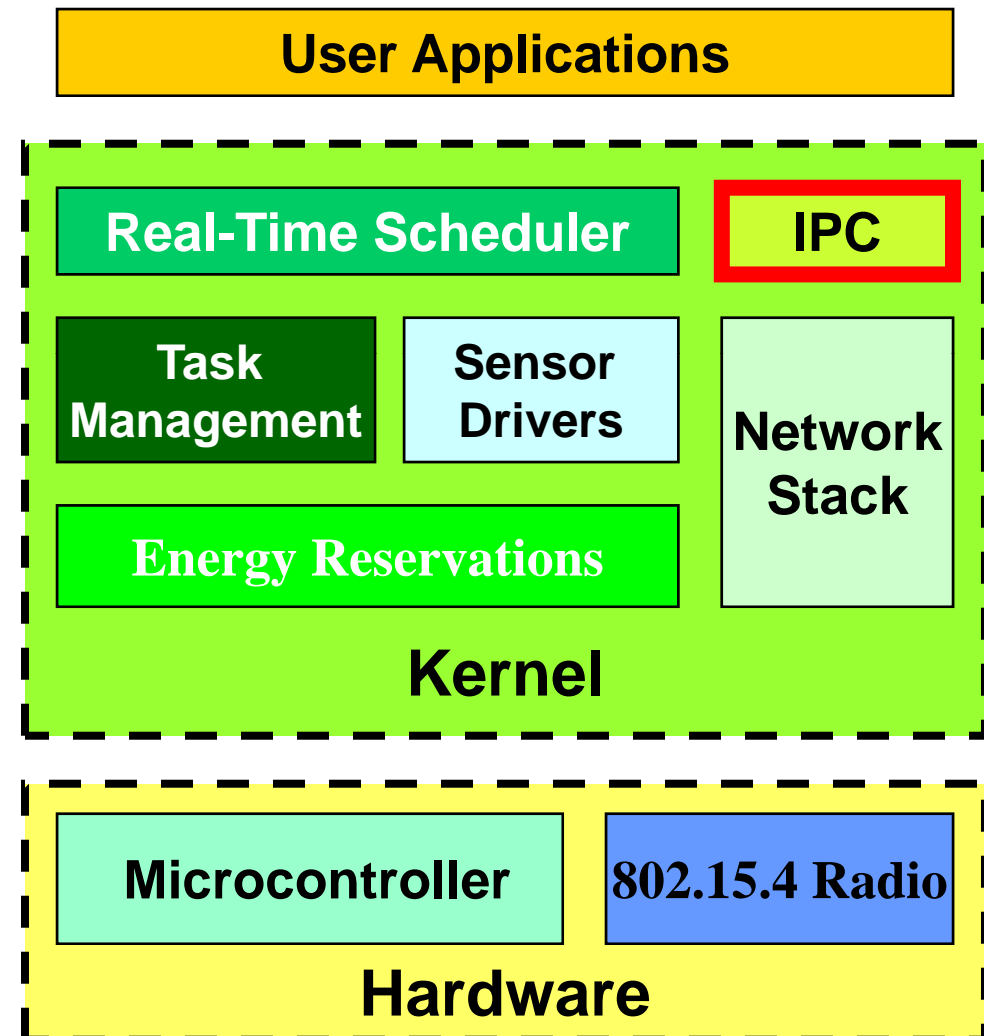
- **System Calls**
 - Allows for arbitration if two tasks read from a non-atomic sensor
- **Raw ADC Values**
 - Temp = 67
- **Real World Values**
 - Temp = 24 (°C)



NanoRK Architecture

- **Inter Process Communication**

- Use Semaphores to arbitrate shared memory communication
- “Message Boxes”



NanoRK Network Stack

- **Zerocopy Buffering Mechanism**
 - TX and RX buffers are in application space
 - Kernel Manages network data through pointer manipulation into application space
- ***Port* Abstraction**
 - A port allows applications to direct their data to individual tasks on each node
- **Network Task**
 - Handles Forwarding Packets, Routing, etc.



Network
Stack

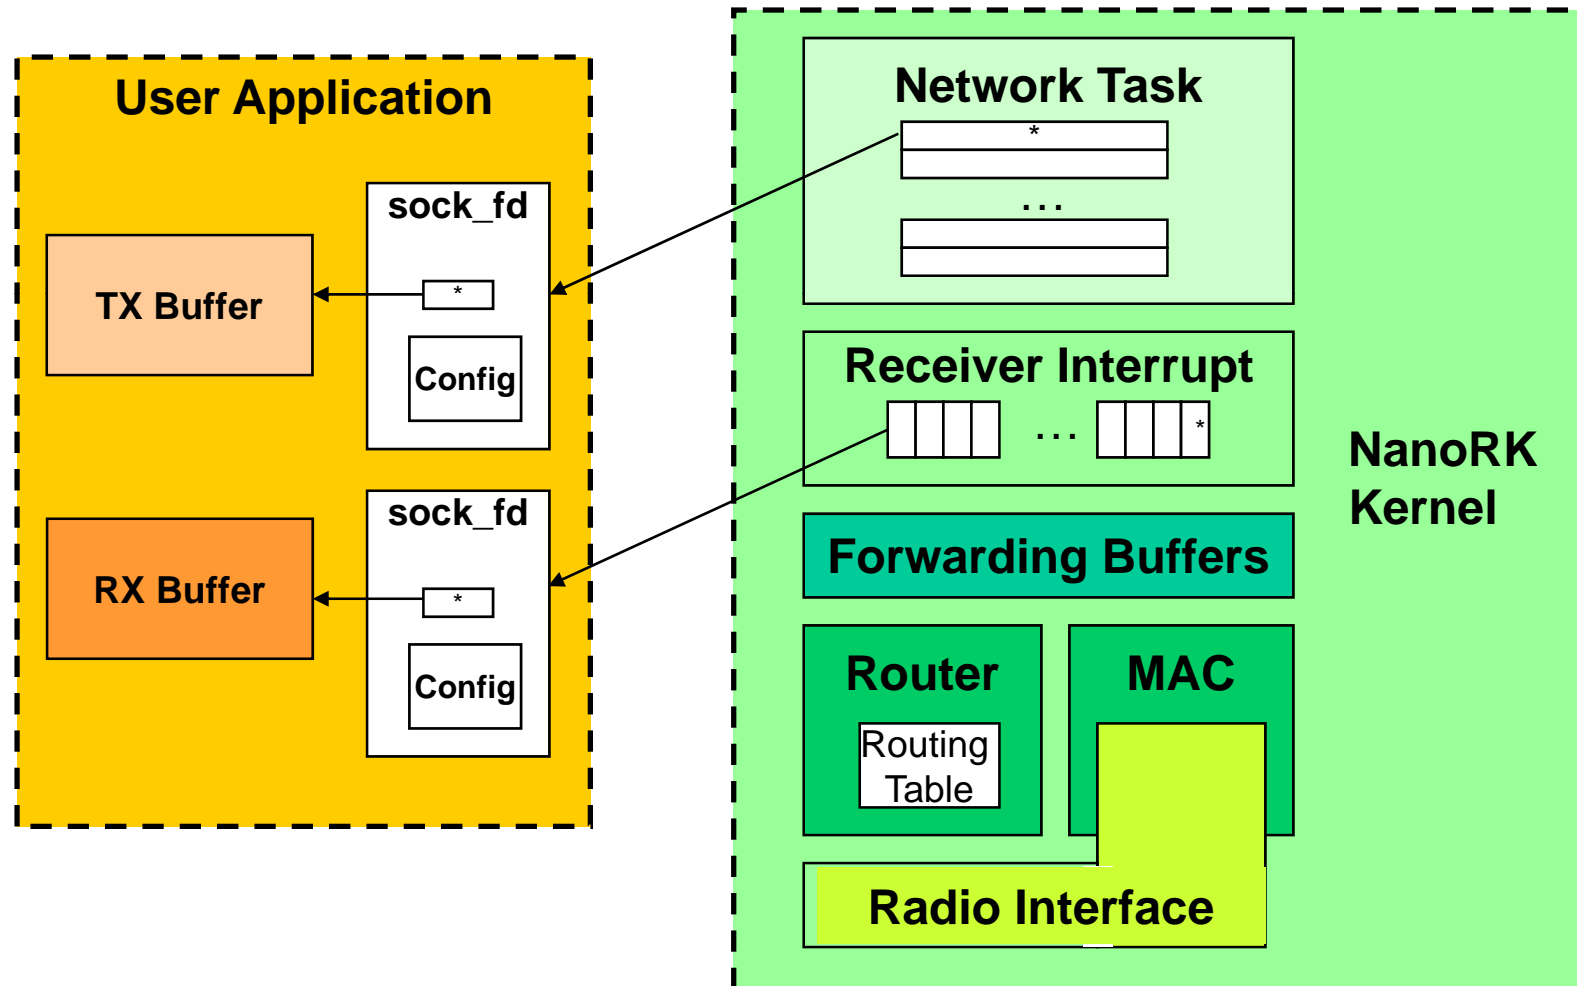
NanoRK Network Stack

- **CSMA MAC Support**
 - Low Power Listen (LPL) MAC support
- **TDMA MAC Support**
 - *RT-Link*: Globally Time Synchronized MAC protocol
- **Ad-Hoc Routing Support**
 - Tasks can access low level routing table
 - Ex: An AODV or DSR TASK can be responsible for establishing and managing routes



Network
Stack

Detailed NanoRK Network Stack



NanoRK Limitations

- **No User / Kernel Space Boundary**
 - Single Memory Space with no MMU
- **Somewhat higher memory overhead**
 - Compared to non-multitasking operating systems

Simulation / Modeling Tool

- **Represent Network as a Graph**

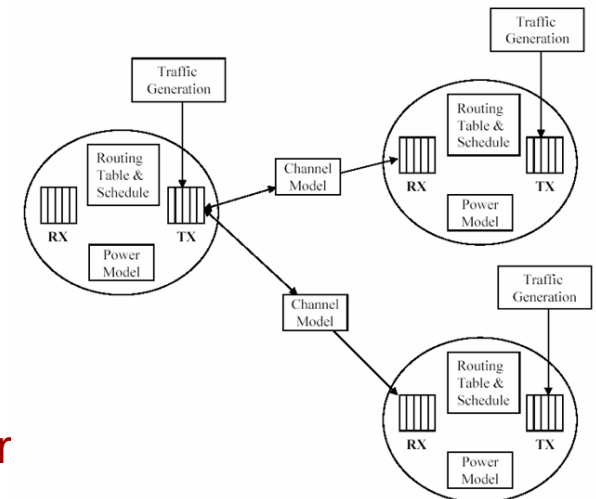
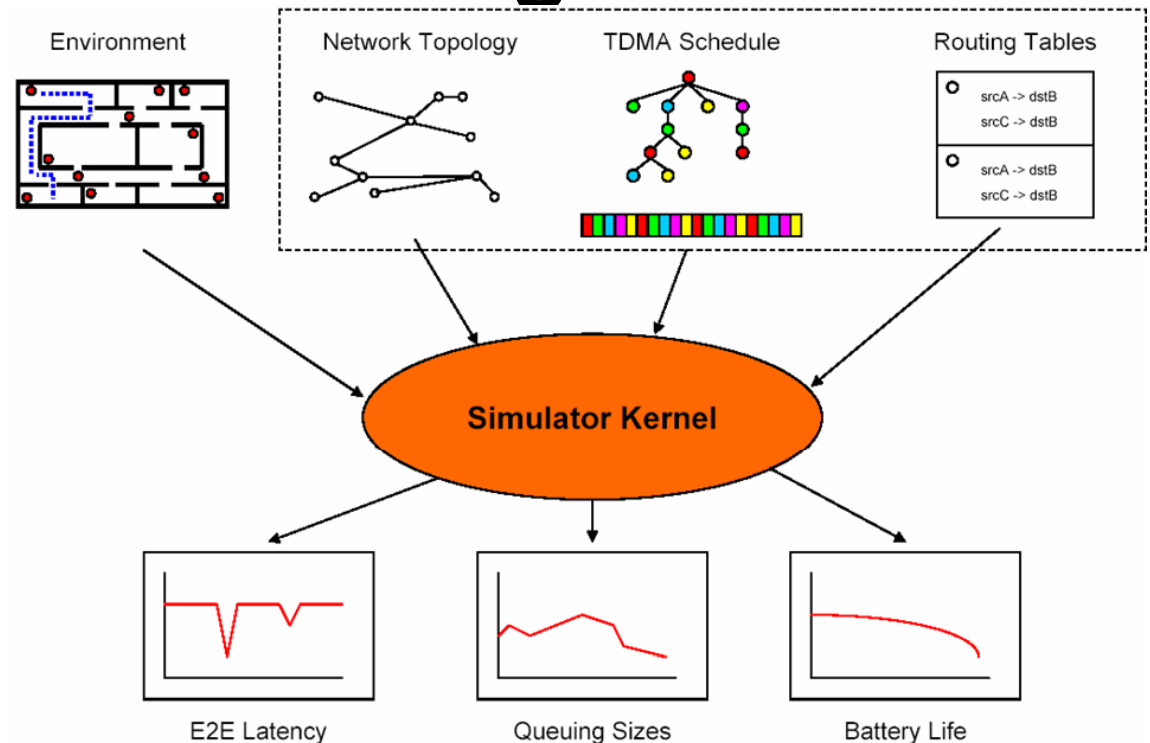
- Easily evaluate different scheduling algorithms

- **Cycle-Accurate Power and Channel Modeling**

- Analyze how different schedules effect power, throughput, and latency

- **Hybrid simulator**

- Allow virtual expansion of our real network
- Works with real nodes



Future Work

- **A public domain release**
- **Dynamic upgrades**
- **Support for multiple microcontrollers**
- **Family of applications**
- **Complete set of tools**
- **Field deployment**

Conclusions

- **NanoRK: a sensor network Operating System**
 - priority-based preemptive multitasking
 - task synchronization
 - multi-hop communications support
- **Reservations can enforce system-wide energy and communication usage**
- **Easy-to-use high level Networking and Sensor Abstractions**

Questions?